Text as Data: What you need to know

Nelson Auner

Prepared for TGG

October 16, 2014

A quick aside...

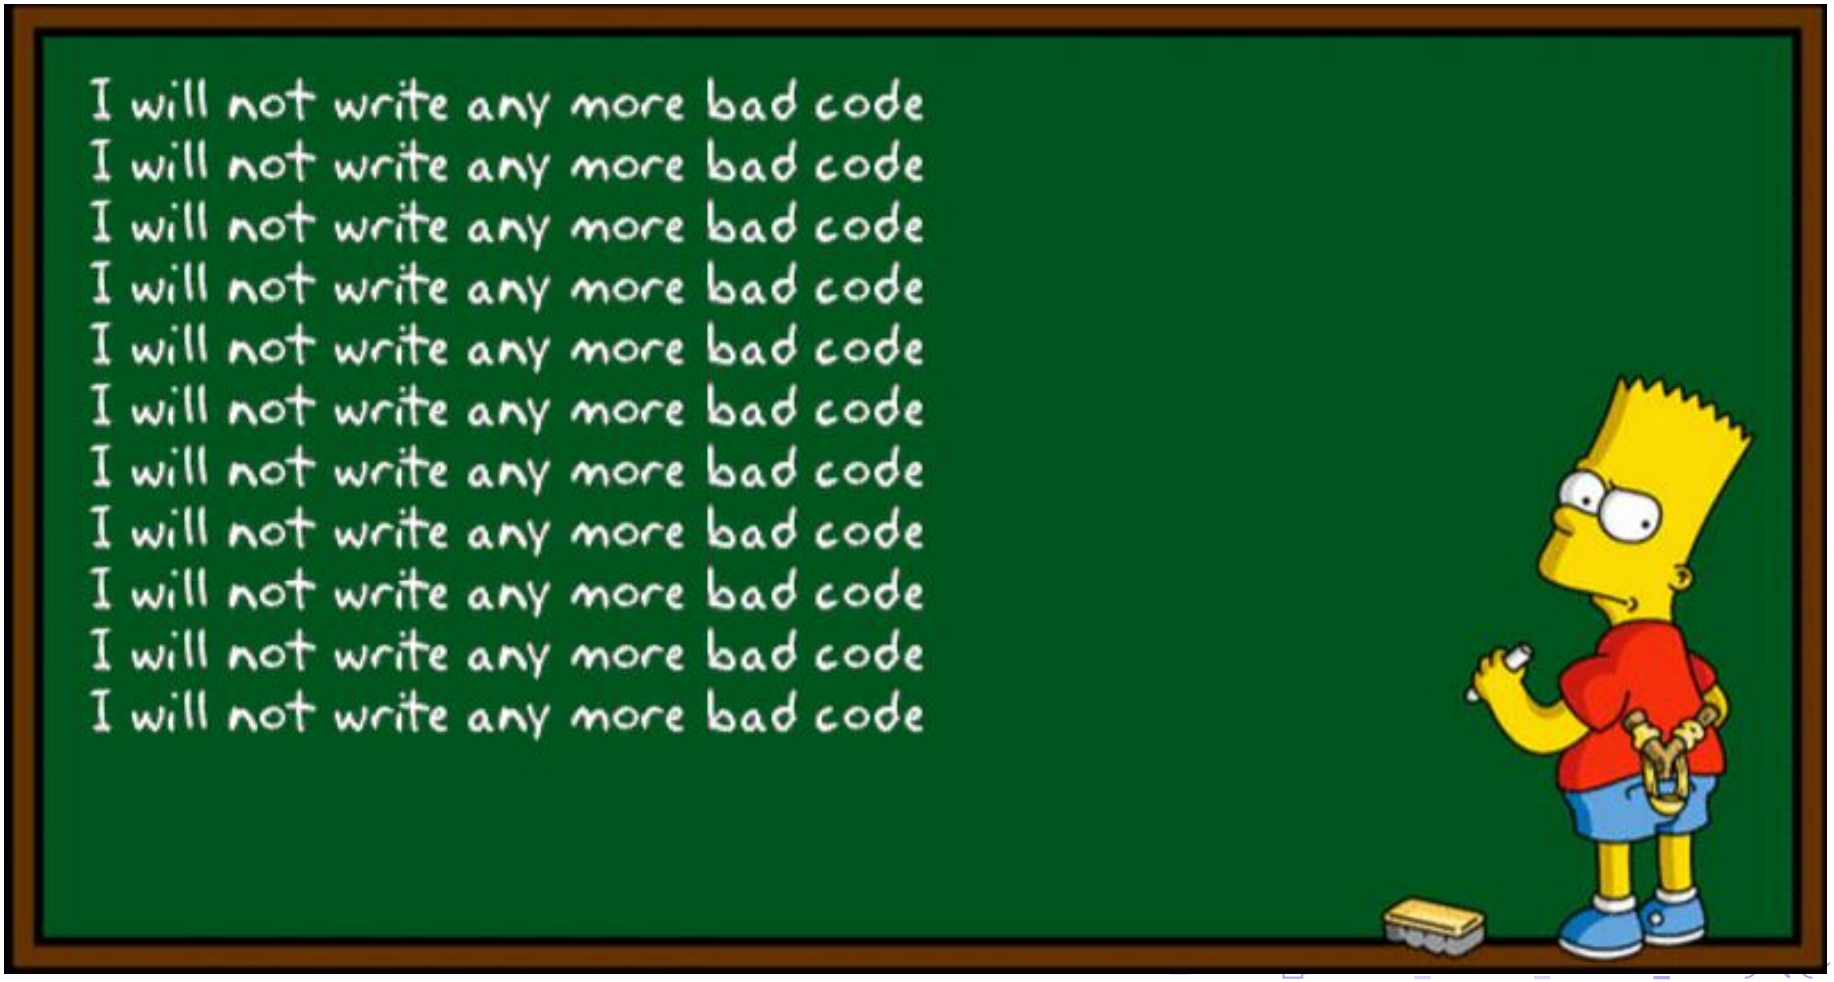

A quick aside...

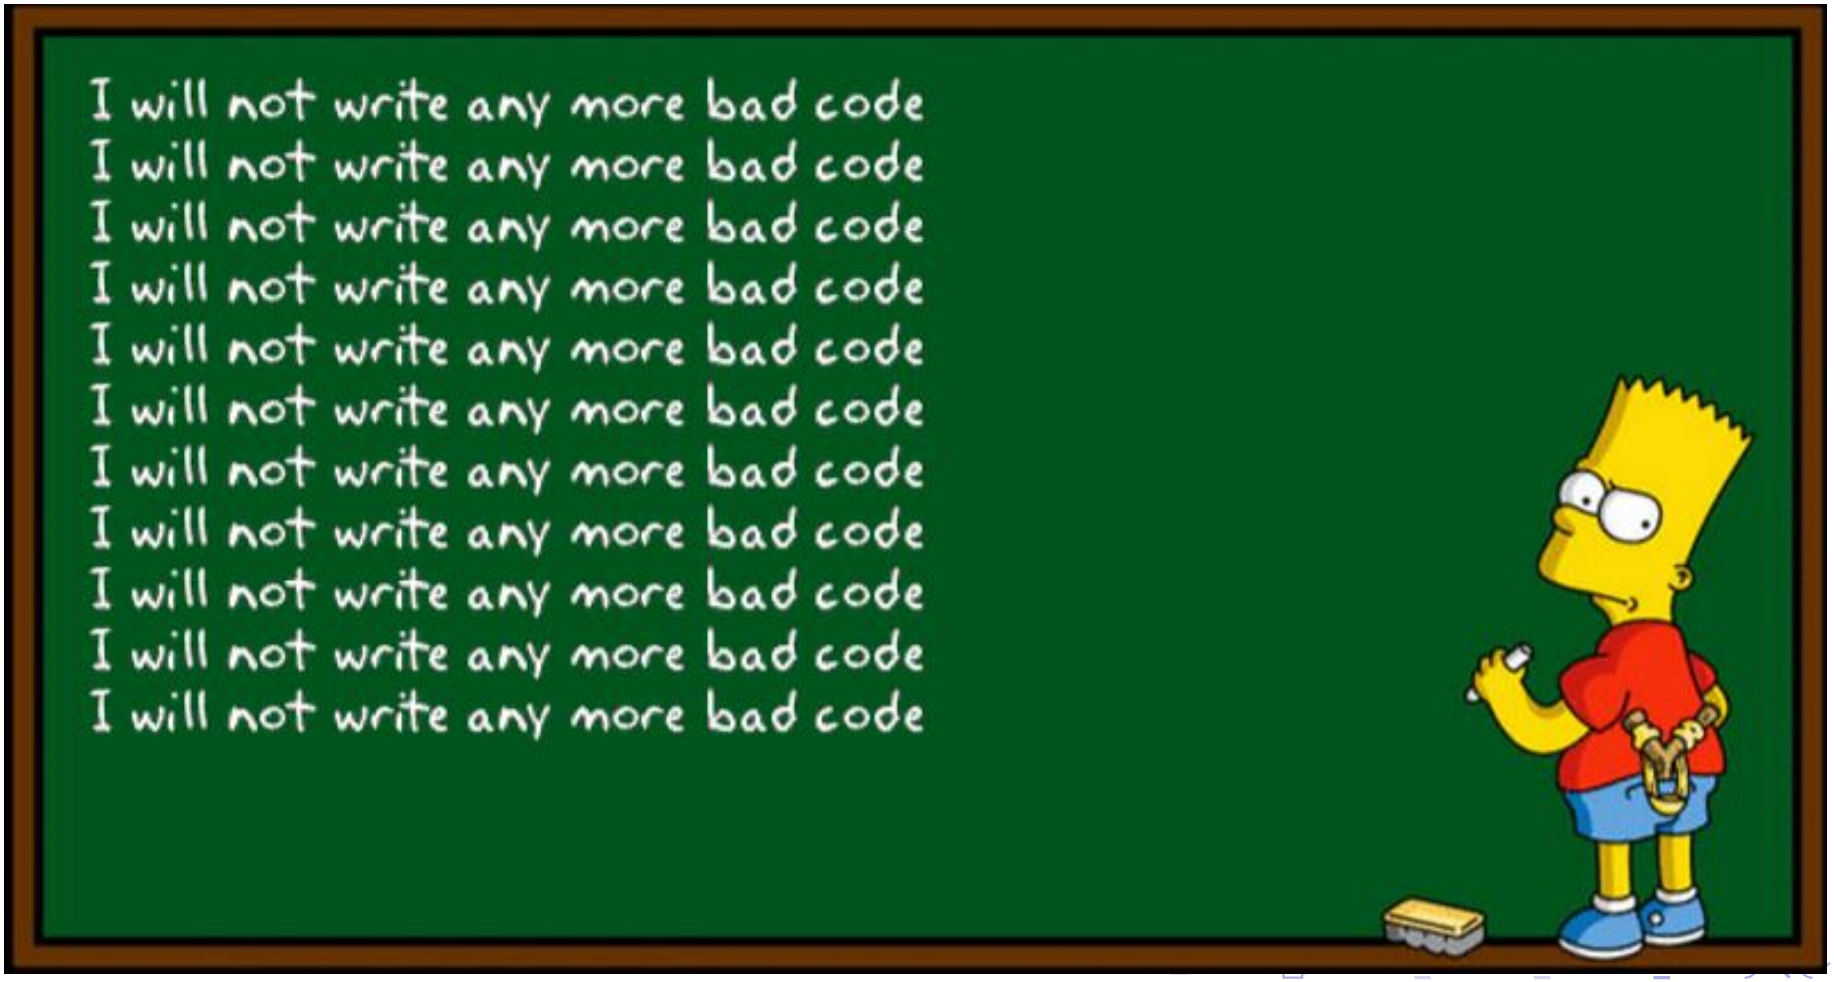

Outline

Motivation: Historical

Motivation: Historical

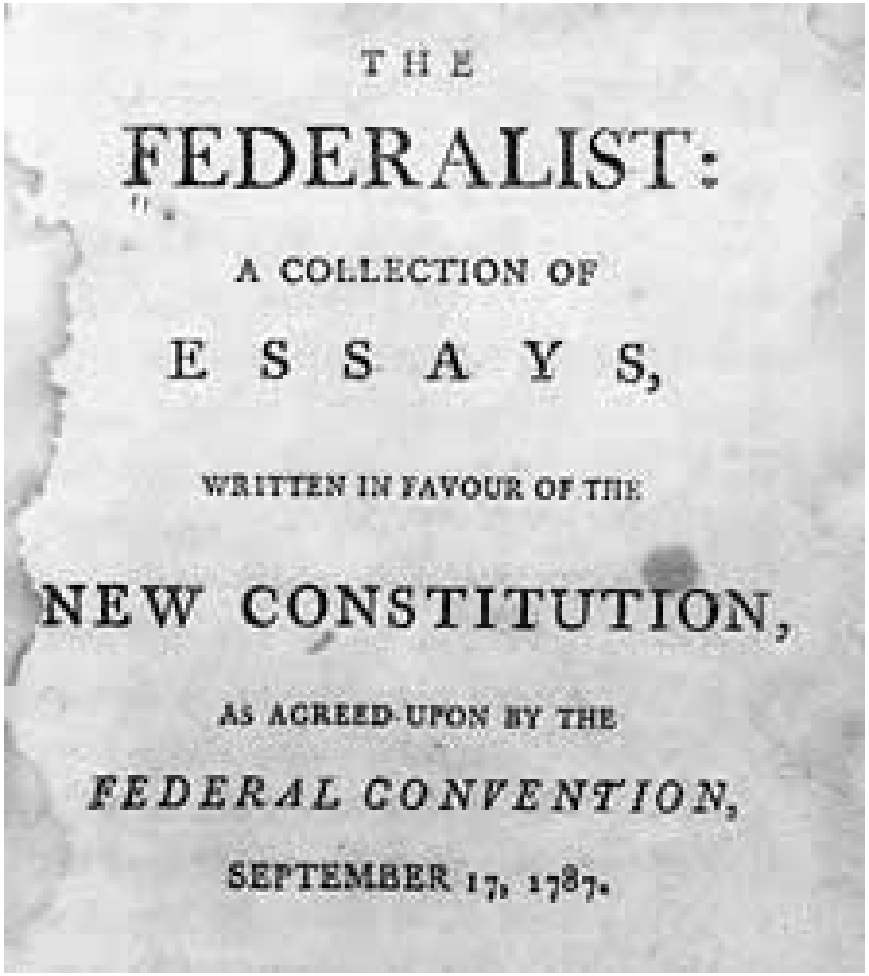
THE
FEDERALIST:
A COLLECTION OF
ESSAYS,
WRITTEN IN FAVOUR OF THE
NEW CONSTITUTION,
AS AGREED UPON BY THE
FEDERAL CONVENTION,
SEPTEMBER 17, 1787.

Motivation: In the News

Motivation: In the News

BloombergBusinessweek
Technology

Global Economics | Companies & Industries | Politics & Policy | Technology | Markets & Finance | Innovation & Design | Lifestyle

Focus On/Big Data

What's the Dollar Value of Online Patient Chatter?

By Caroline Chen | October 09, 2014

SEND TO kindle

On the Internet, people talk about everything, even their illnesses. Treato, an Israeli data-mining company, monitors conversations on Facebook (FB), Twitter (TWTR), and patient forums for information on drug side effects and prescription patterns. It then sends weekly or monthly analyses of the chatter to hedge funds and money managers that invest in pharma stocks.

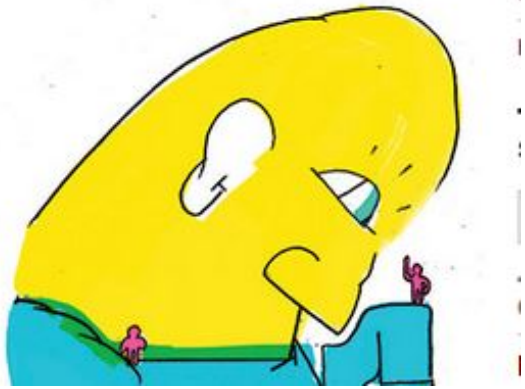

How do they do it?

How do they do it?

- "Treato distills the collective patient voice from blogs and forums using Natural Language Processing, Big Data and a proprietary patient language..."

Public Sector Use

Public Sector Use

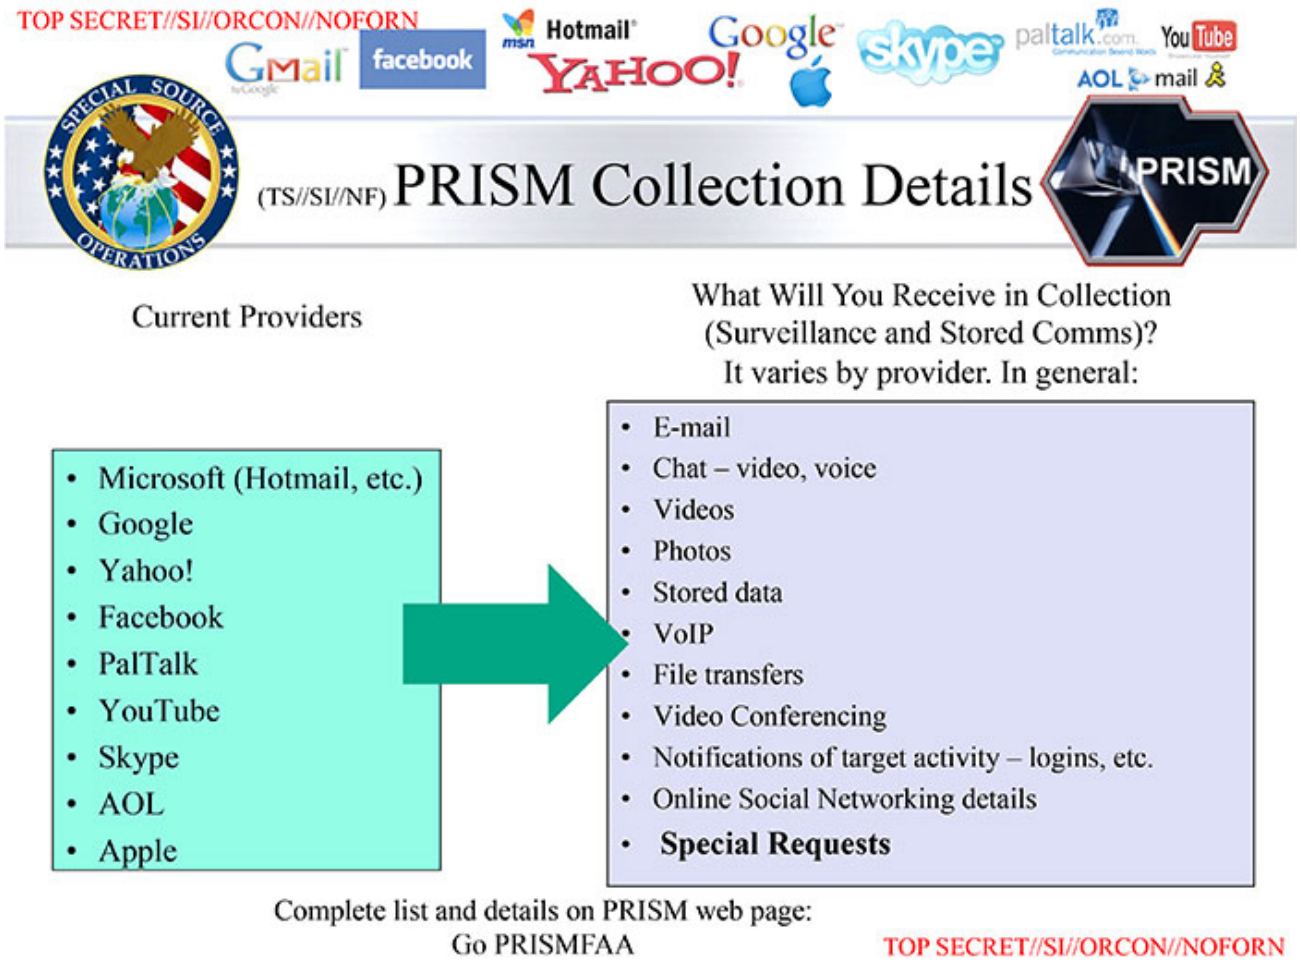

Outline

Are you comfortable talking about text?

Are you comfortable talking about text?

Can you...

Are you comfortable talking about text?

Can you...

- explain the basics of text analysis to a potential client?

Are you comfortable talking about text?

Can you...

- explain the basics of text analysis to a potential client?
- identify opportunity to utilize text analysis?

Are you comfortable talking about text?

Can you...

- explain the basics of text analysis to a potential client?
- identify opportunity to utilize text analysis?
- commicate why text analysis is difficult?

Outline

The basics

The basics

- A document is a collection of words or phrases.

The basics

- A document is a collection of words or phrases.
- Our datasets are collections of documents

The basics

- A document is a collection of words or phrases.
- Our datasets are collections of documents

Table: What did homework consist of?

The basics

- A document is a collection of words or phrases.
- Our datasets are collections of documents

Table: What did homework consist of?

Document	Content
1	Some computation and formula proving, a lot of R code
2	Problems, computation using R
3	Some computations and writing R code
4	Proofs, problems, and programming work

Parsing

Parsing

- Greatest, Greatly, and Greatliest....

Parsing

- Greatest, Greatly, and Greatliest....

It ain't that easy...

Parsing

- Greatest, Greatly, and Greatliest....

It ain't that easy...

Crystial rosey yeah I poe that
We connected with Cali we back door that

Parsing

- Greatest, Greatly, and Greatliest....

It ain't that easy...

Crystial rosey yeah I poe that
We connected with Cali we back door that
You see my wrist man keep your pink wrist bands
She can't believe I'm in a chevy even though I'm rich man

Parsing

- Greatest, Greatly, and Greatliest....

It ain't that easy...

Crystial rosey yeah I poe that
We connected with Cali we back door that
You see my wrist man keep your pink wrist bands
She can't believe I'm in a chevy even though I'm rich man
Chevy Ridin' High - Dre (of Cool and Dre) f/ Rick Ross

Mo'(Multinomial) Models

Mo'(Multinomial) Models

- If word order doesn't matter, then we can treat each document as a "bag of words".

Mo'(Multinomial) Models

- If word order doesn't matter, then we can treat each document as a "bag of words".
- The number of words can be modeled $\sim$ multinomial

Mo'(Multinomial) Models

- If word order doesn't matter, then we can treat each document as a "bag of words".
- The number of words can be modeled $\sim$ multinomial

Table: Creating a word-count matrix from text

Document	Some	comp	formula	prov	R	code	use	problem	writ	program	work
1	1	1	1	1	1	1	0	0	0	0	0
2	0	1	0	0	1	0	1	1	0	0	0
3	1	1	0	0	1	0	0	0	1	0	0
4	0	0	0	1	0	0	0	1	0	1	1

A better model: Metadata

- We would like to add structure to the model for inference or prediction

A better model: Metadata

- We would like to add structure to the model for inference or prediction
- Metadata is data that accompanies a document

A better model: Metadata

- We would like to add structure to the model for inference or prediction
- Metadata is data that accompanies a document

Table: What did homework consist of?

A better model: Metadata

- We would like to add structure to the model for inference or prediction
- Metadata is data that accompanies a document

Table: What did homework consist of?

Grade	**Content**
A+	Some computation and formula proving, a lot of R code
B	Problems, computation using R
B	Some computations and writing R code
C+	Proofs, problems, and programming work

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

[1]Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2]Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

[1]Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2]Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

Bike Topic

Pedal, Helmet, Gears

[1] Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2] Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

Bike Topic

Pedal, Helmet, Gears

Swimming

Stroke, Air, Water

[1] Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2] Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

Bike Topic

Pedal, Helmet, Gears

Swimming

Stroke, Air, Water

- A book about triathalon training $\sim \theta_1$ Running $+ \theta_2$ Biking $+$ θ_3 Swimming

[1] Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2] Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

Bike Topic

Pedal, Helmet, Gears

Swimming

Stroke, Air, Water

- A book about triathalon training $\sim \theta_1$ Running $+ \theta_2$ Biking $+ \theta_3$ Swimming
- Issue: We can no longer collapse observations, must use all n observations

[1] Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2] Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Topic Models

A topic is a distribution of words.

In a topic model, documents are made of a mixtures of topics.

Running Topic

Stride, Pacing, Stretch

Bike Topic

Pedal, Helmet, Gears

Swimming

Stroke, Air, Water

- A book about triathalon training $\sim \theta_1$ Running $+ \theta_2$ Biking $+ \theta_3$ Swimming
- Issue: We can no longer collapse observations, must use all n observations
- Workarounds: See Ryan's paper[1] or mine [2]

[1] Wang, 2012. Sparse Coding and an Application to Topic Modeling.

[2] Auner, 2014. Combining Latent Topics with Document Attributes in Text Analysis

Outline

Cluster Model

Goal

- Want to use the Topic Model but incorporate Metadata
- Also want computational ease

Cluster Model

Goal

- Want to use the Topic Model but incorporate Metadata
- Also want computational ease

Approach

- Restrict each document to only one topic $\Rightarrow$ "cluster"
- Can collapse observations over unique (metadata, cluster) combination
- $x_i \sim MN(q_{ij}, m_{ij})$; $\quad q_{ij} = \frac{exp(\alpha_j + y_i\phi_j + u_i\Gamma_{kj})}{\sum_{l=1}^{p} exp(\alpha_l + y_i\phi_l + u_i\Gamma_{kl})}$

Algorithm for Cluster Membership Model with Gamma Lasso Penalty

1. Initialize cluster membership u_i for $i = 1, \ldots, n$

Algorithm for Cluster Membership Model with Gamma Lasso Penalty

1. Initialize cluster membership u_i for $i = 1, \ldots, n$
2. Determine parameters α, ϕ, Γ by fitting a multinomial regression on $y_i | x_i, u_i$ with a gamma lasso penalty (Taddy 2013)

Algorithm for Cluster Membership Model with Gamma Lasso Penalty

1. Initialize cluster membership u_i for $i = 1, \ldots, n$
2. Determine parameters α, ϕ, Γ by fitting a multinomial regression on $y_i | x_i, u_i$ with a gamma lasso penalty (Taddy 2013)
3. For each document i, determine new cluster u_i membership as $argmax_{k=1,..,K} \left[\ell(u_i | \alpha, \phi, \Gamma) \right]$

Algorithm for Cluster Membership Model with Gamma Lasso Penalty

1. Initialize cluster membership u_i for $i = 1, \ldots, n$
2. Determine parameters α, ϕ, Γ by fitting a multinomial regression on $y_i|x_i, u_i$ with a gamma lasso penalty (Taddy 2013)
3. For each document i, determine new cluster u_i membership as $argmax_{k=1,..,K}\left[\ell(u_i|\alpha, \phi, \Gamma)\right]$
4. Check if current cluster assignment is different from previous cluster assignment , $(\mathbf{u}^{(t)} = \mathbf{u}^{(t-1)})$.If so, return to step 2. If not, end algorithm.

Outline

Comparison with the Topic Model

Good news: We are able to recover similar topics with our model:

Table: Comparison of top word loadings on a stem-cell topic

Cluster Membership	Topic Model (LDA)*
umbilic.cord.blood	pluripotent.stem.cel
cord.blood.stem	national.ad.campaign
blood.stem.cel	cel.stem.cel
adult.stem.cel	stem.cel.line

*Results reported in Taddy (2012)

Incorporating metadata: Restaurant Review

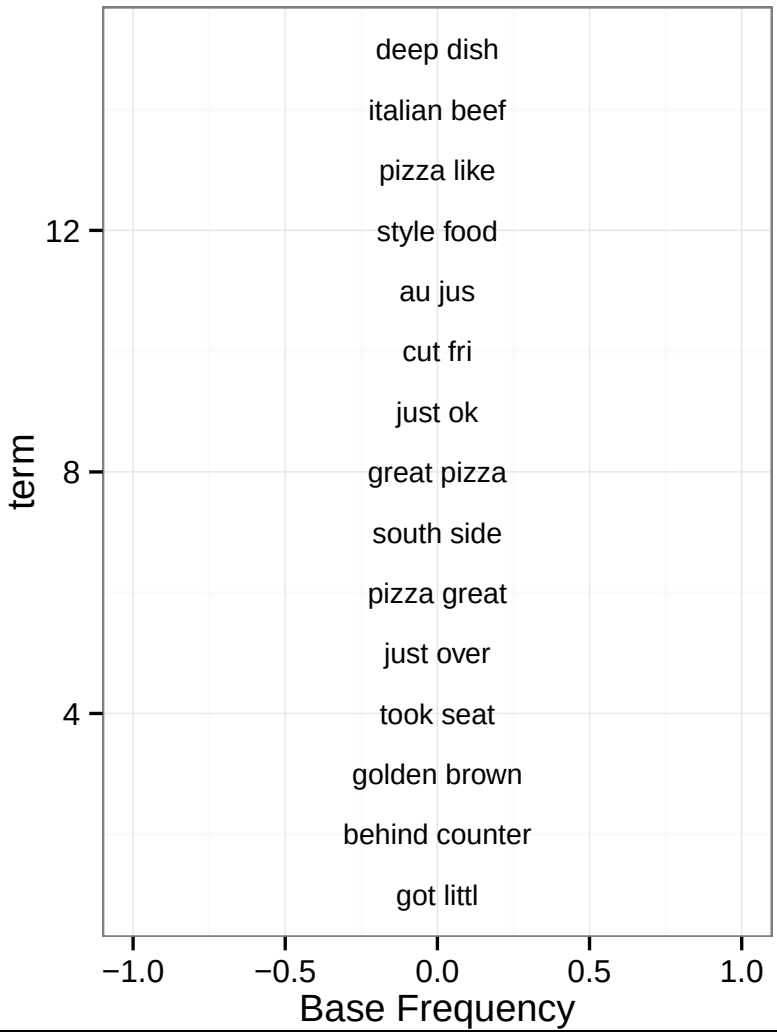

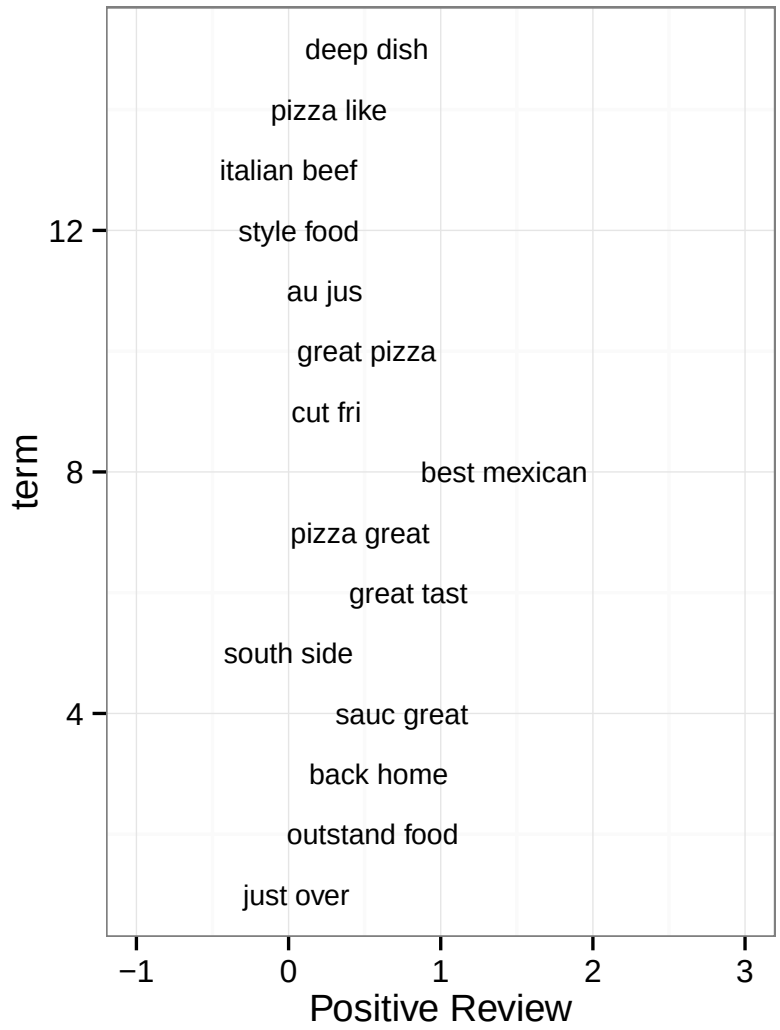

Imma Let you Finish, but the Dirichlet was the greatest prior of all time!

Imma Let you Finish, but the Dirichlet was the greatest prior of all time!

Results

Results

	term	loading
1	yeezus	5.48
2	constel	3.79
3	homm	3.79
4	preach	3.79
5	bound	3.6
6	thoma	3.38
7	thirti	3.32
8	rocka	3.31
9	rowland	3.25
10	jamaican	3.23
11	blocka	3.22
12	movement	3.22
13	unlik	3.08

Dip your feet in

Dip your feet in

- *Textir* or *Gamlr* package by Matt Taddy
- Currently only for R
- Python coming soon!

Thank You

Thank You