


This article was downloaded by: [University of Chicago Library]

On: 14 May 2014, At: 10:23

Publisher: Taylor & Francis

Informa Ltd Registered in England and Wales Registered Number: 1072954 Registered office: Mortimer House, 37-41 Mortimer Street, London W1T 3JH, UK



Technometrics

Publication details, including instructions for authors and subscription information:

<http://www.tandfonline.com/loi/utch20>

Measuring Political Sentiment on Twitter: Factor Optimal Design for Multinomial Inverse Regression

Matt Taddy^a

^a The University of Chicago Booth School of Business, Chicago, IL, 60637

Accepted author version posted online: 07 Mar 2013. Published online: 22 Nov 2013.

To cite this article: Matt Taddy (2013) Measuring Political Sentiment on Twitter: Factor Optimal Design for Multinomial Inverse Regression, *Technometrics*, 55:4, 415-425, DOI: [10.1080/00401706.2013.778791](https://doi.org/10.1080/00401706.2013.778791)

To link to this article: <http://dx.doi.org/10.1080/00401706.2013.778791>

PLEASE SCROLL DOWN FOR ARTICLE

Taylor & Francis makes every effort to ensure the accuracy of all the information (the "Content") contained in the publications on our platform. However, Taylor & Francis, our agents, and our licensors make no representations or warranties whatsoever as to the accuracy, completeness, or suitability for any purpose of the Content. Any opinions and views expressed in this publication are the opinions and views of the authors, and are not the views of or endorsed by Taylor & Francis. The accuracy of the Content should not be relied upon and should be independently verified with primary sources of information. Taylor and Francis shall not be liable for any losses, actions, claims, proceedings, demands, costs, expenses, damages, and other liabilities whatsoever or howsoever caused arising directly or indirectly in connection with, in relation to or arising out of the use of the Content.

This article may be used for research, teaching, and private study purposes. Any substantial or systematic reproduction, redistribution, reselling, loan, sub-licensing, systematic supply, or distribution in any form to anyone is expressly forbidden. Terms & Conditions of access and use can be found at <http://www.tandfonline.com/page/terms-and-conditions>

Measuring Political Sentiment on Twitter: Factor Optimal Design for Multinomial Inverse Regression

Matt TADDY

The University of Chicago Booth School of Business
Chicago, IL 60637
(taddy@chicagobooth.edu)

This article presents a short case study in text analysis: the scoring of Twitter posts for positive, negative, or neutral sentiment directed toward particular U.S. politicians. The study requires selection of a subsample of representative posts for sentiment scoring, a common and costly aspect of sentiment mining. As a general contribution, our application is preceded by a proposed algorithm for maximizing sampling efficiency. In particular, we outline and illustrate greedy selection of documents to build designs that are D-optimal in a topic-factor decomposition of the original text. The strategy is applied to our motivating dataset of political posts, and we outline a new technique for predicting both generic and subject-specific document sentiment through the use of variable interactions in multinomial inverse regression. Results are presented for analysis of 2.1 million Twitter posts collected around February 2012. Computer codes and data are provided as supplementary material online.

KEY WORDS: Active learning; Multinomial logistic regression; Sentiment mining; Text analysis; Topic models; Variable interaction.

1. INTRODUCTION

This article outlines a simple approach to a general problem in text analysis, the selection of documents for costly annotation. We then show how inverse regression can be applied with variable interactions to obtain both generic and subject-specific predictions of document sentiment, our annotation of interest. We are motivated by the problem of design and analysis of a particular text mining experiment: the scoring of Twitter posts (“tweets”) for positive, negative, or neutral sentiment directed toward particular U.S. politicians. The contribution is structured first with a proposal for optimal design of text data experiments, followed by application of this technique in our political tweet case study and analysis of the resulting data through inverse regression.

Text data are viewed throughout simply as counts, for each document, of phrase occurrences. These phrases can be words (e.g., *tax*) or word combinations (e.g., *pay tax* or *too much tax*). Although there are many different ways to process raw text into these *tokens*, perhaps using sophisticated syntactic or semantic rules, we do not consider the issue in detail and assume tokenization as given; our case study text processing follows a few simple rules described below. Document i is represented as $\mathbf{x}_i = [x_{i1}, \dots, x_{ip}]'$, a sparse vector of counts for each of p tokens in the vocabulary, and a document-term count matrix is written $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]'$, where n is the number of documents in a given corpus. These counts, and the associated frequencies $\mathbf{f}_i = \mathbf{x}_i / m_i$ where $m_i = \sum_{j=1}^p x_{ij}$, are then the basic data units for statistical text analysis. Hence, text data can be characterized simply as exchangeable counts in a very large number of categories, leading to the common assumption of a multinomial distribution for each \mathbf{x}_i .

We are concerned with predicting the *sentiment* $\mathbf{y} = [y_1, \dots, y_n]'$ associated with documents in a corpus. In our main application, this is positive, neutral, or negative sentiment directed toward a given politician, as measured through a reader survey. More generally, sentiment can be replaced by any annotation that is correlated with document text. Text-sentiment prediction is thus just a very high-dimensional regression problem, where the covariates have the special property that they can be represented as draws from a multinomial distribution.

Any regression model needs to be accompanied with data for training. In the context of sentiment prediction, this implies documents scored for sentiment. One can look to various sources of “automatic” scoring, and these are useful to obtain the massive amounts of data necessary to train high-dimensional text models. Section 1.1 describes our use of emoticons for this purpose. However, such automatic scores are often only a rough substitute for the true sentiment of interest. In our case, generic happy/sad sentiment is not the same as sentiment directed toward a particular politician. It is then necessary to have a subset of the documents annotated with precise scores, and since this scoring will cost money, we need to choose a subset of documents whose content is most useful for predicting sentiment from text. This is an application for *pool-based active learning*: there is a finite set of examples for which predictions are to be obtained and one seeks to choose an optimal representative subset.

There are thus two main elements to our study: design, choosing the subsample of tweets to be sent for scoring, and analysis, using sentiment-scored tweets to fit a model for predicting Twitter sentiment toward specific politicians. This article is about both components. As a design problem, text mining presents a difficult situation where raw space filling is impractical—the dimension of \mathbf{x} is so large that every document is very far apart—and we argue in Section 3 that it is unwise to base design choices on the poor estimates of predictive uncertainty provided by text regression. Our solution is to use a space-filling design but in an estimated lower dimensional multinomial-factor space rather than in the original \mathbf{x} -sample. Section 3.1 describes a standard class of *topic models* that can be used to obtain low-dimensional factor representations for large document collections. The resulting unsupervised algorithm (i.e., sampling proceeds without regard to sentiment) can be combined with any sentiment prediction model. We use the multinomial inverse regression (MNIR) of Taddy (2013), with the addition of politician-specific interaction terms, as described in Section 2.

1.1 Data Application: Political Sentiment on Twitter

The motivating case study for this article is an analysis of sentiment in tweets about U.S. politicians on Twitter, the social blog, from January 27 to February 28, 2012, a period that included the Florida (1/31), Nevada (2/4), Colorado, Missouri, and Minnesota (2/7), Maine (2/11), and Michigan and Arizona (2/28) presidential primary elections. Twitter provides streaming access to a large subset of public (as set by the user) tweets containing terms in a short list of case insensitive filters. We were interested in conversation on the leading candidates in the Republican presidential primary, as well as that concerning current president Barack Obama; our list of filter terms was *obama*, *romney*, *gingrich*, *ron paul*, and, from February 13 onward, *santorum*. Note that Romney, Gingrich, and Paul were the only front-runners at the beginning of our study, but Santorum gained rapidly in the polls following his surprise victories in three state votes on February 7: the Minnesota and Colorado caucuses and the Missouri Primary. Daily data collection is shown by politician-subject in Figure 1; total counts are 10.2×10^5 for Obama, 5×10^5 for

Romney, 2.2×10^5 for Gingrich, 2.1×10^5 for Santorum, and 1.5×10^5 for Paul, for a full sample of about 2.1 million tweets.

In processing the raw text, we remove a limited set of stop words (terms that occur at a constant rate regardless of subject, such as *and* or *the*) and punctuation before converting to lowercase and stripping suffixes from roots according to the Porter stemmer (Porter 1980). The results are then tokenized into single terms based upon separating white-space, and we discard any tokens that occur in <200 tweets and are not in the list of tokens common in our generic emoticon-sentiment tweets, described in the next paragraph. This leads to 5532 unique tokens for Obama, 5352 for Romney, 5143 for Gingrich, 5131 for Santorum, and 5071 for Paul.

The primary analysis goal is to classify tweets by sentiment: positive, negative, or neutral. We have two data sources available: twitter data that are scored for generic sentiment and the ability to survey readers about sentiment in tweets directed at specific politicians. In the first case, 1.6 million tweets were obtained, from the website <http://twittersentiment.appspot.com>, that have been automatically identified as positive or negative by the presence of an emoticon (symbols included by the author—for example, a happy face indicates a positive tweet and a sad face a negative tweet). Tokenization for these tweets followed the same rules as for the political Twitter sample above, and we discard tokens that occur in less than 0.01% of tweets. This leads to a vocabulary of 5412 “emoticon” tokens; due to considerable overlap, the combined vocabulary across all tweets (political and emoticon) is only 5690 tokens.

As our second data source, we use the Amazon Mechanical Turk (<https://www.mturk.com/>) platform for scoring tweet sentiment. Tweets are shown to anonymous workers for categorization as representing either positive (e.g., “support, excitement, respect, or optimism”) or negative (e.g., “anger, distrust, disapproval, or ridicule”) feelings or news toward a given politician, or as neutral if the text is “irrelevant, or not even slightly positive or negative.” Each tweet is seen by two independent workers, and it is only considered scored if the two agree on categorization. In addition, workers were prescreened as “masters” by Amazon and we monitored submissions for quality control, blocking poor workers. Given the 2–3 cents per tweet paid to individual workers, as well as the overhead charged by Amazon,




Figure 1. Tweet sample volume for political candidates. All are taken from the stream of public Twitter posts from January 27 through the end of February, except for Santorum who was only tracked after February 13. The online version of this figure is in color.

our worker agreement rates of around 80% imply an average cost near \$0.075 per sentiment scored tweet.

2. SENTIMENT PREDICTION VIA MULTINOMIAL INVERSE REGRESSION

Sentiment prediction in this article follows the MNIR framework described in Taddy (2013). Section 2.1 summarizes that approach, while Section 2.2 discusses an adaptation specific to the main application of this article. Inverse regression as a general strategy looks to estimate the *inverse distribution* for covariates given response and to use this as a tool in building a *forward model* for y_i given \mathbf{x}_i . The specific idea of MNIR is to estimate a simple model for how the multinomial distribution on text counts changes with sentiment and to derive from this model low-dimensional text projections that can be used for predicting sentiment.

2.1 Single-Factor MNIR

As a simple case, suppose that y_i for document i is a discrete ordered sentiment variable with support \mathcal{Y} —say $y_i \in \{-1, 0, 1\}$ as in our motivating application. Only a very complicated model will be able to capture the generative process for an individual's text, $\mathbf{x}_i|y_i$, which involves both heterogeneity between individuals and correlation across dimensions of \mathbf{x}_i . Thus estimating a model for $\mathbf{x}_i|y_i$ can be far harder than predicting y_i from \mathbf{x}_i , and inverse regression does not seem a clever place to be starting analysis. However, we can instead concentrate on the *population average* effect of sentiment on text by modeling the conditional distribution for collapsed token counts $\mathbf{x}_y = \sum_{i:y_i=y} \mathbf{x}_i$. A basic MNIR model is then

$$\mathbf{x}_y \sim \text{MN}(\mathbf{q}_y, m_y) \quad \text{with} \quad q_{yj} = \frac{\exp[\alpha_j + y\varphi_j]}{\sum_{l=1}^p \exp[\alpha_l + y\varphi_l]}, \quad (1)$$

for $j = 1, \dots, p, y \in \mathcal{Y}$,

where each MN is a p -dimensional multinomial distribution with size $m_y = \sum_{i:y_i=y} m_i$ and probabilities $\mathbf{q}_y = [q_{y1}, \dots, q_{yp}]'$ that are a linear function of y through a logistic link. Although independence assumptions implied by (1) are surely incorrect, within-individual correlation in \mathbf{x}_i is quickly overwhelmed in aggregation and the multinomial becomes decent model for \mathbf{x}_y . (One could also argue against an equidistant three point scale for y , however such a scale is useful to simplify inverse regression and we assume that misspecification here can be accommodated in forward regression.)

Given sentiment y and counts \mathbf{x} drawn from the multinomial distribution $\text{MN}(\mathbf{q}_y, m)$ in (1), the projection $\boldsymbol{\varphi}'\mathbf{x}$ is *sufficient for sentiment* in the sense that $y \perp\!\!\!\perp \mathbf{x} \mid \boldsymbol{\varphi}'\mathbf{x}, m$. A simple way to demonstrate this is through application of Bayes rule (after assigning prior probabilities for each element of \mathcal{Y}). Then given \mathbf{x}_i counts for an *individual* document, $\boldsymbol{\varphi}'\mathbf{x}_i$ seems potentially useful as a low-dimensional index for predicting y_i . More specifically, we normalize by document length in defining the *sufficient reduction* (SR) score

$$z_i = \boldsymbol{\varphi}'\mathbf{f}_i = \boldsymbol{\varphi}'\mathbf{x}_i / m_i. \quad (2)$$

Now, since (1) is a model for collapsed text counts rather than for \mathbf{x}_i given y_i , the SR score in (2) is *not* theoretically sufficient

for that document's sentiment. Taddy (2013) described specific random effects models for the information loss in regressing y_i onto z_i instead of \mathbf{x}_i , and under certain models, the individual document regression coefficients approach $\boldsymbol{\varphi}$. However, in general, this population average projection is *misspecified* as an individual document projection. Hence, instead of applying Bayes rule to invert (1) for sentiment prediction, z_i is treated as observable in a second-stage regression for y_i given z_i . Throughout this article, where y is always an ordered discrete sentiment variable, this *forward regression* applies logistic proportional odds models of the form $p(y_i < c) = (1 + \exp[-(\gamma_c + \beta z_i)])^{-1}$.

2.2 MNIR With Politician-Interaction

In the political twitter application, our approach needs to be adapted to allow different text-sentiment regression models for each politician and also to accommodate positive and negative emoticon tweets, which are sampled from all public tweets rather than always being associated with a politician. This is achieved naturally within the MNIR framework by introducing interaction terms in the inverse regression.

The data are now written with text in the i th tweet for politician s as \mathbf{x}_{si} , containing a total of m_{si} tokens and accompanied by sentiment $y_{si} \in \{-1, 0, 1\}$, corresponding to negative, neutral, and positive sentiment, respectively. Collapsed counts for each politician-sentiment combination are obtained as $x_{syj} = \sum_{i:y_{si}=y} x_{sij}$ for each token j . This yields 17 “observations”: each of three sentiments for five politicians, plus positive and negative emoticon tweets. The MNIR model for sentiment- y text counts directed toward politician s is then $\mathbf{x}_{sy} \sim \text{MN}(\mathbf{q}_{sy}, m_{sy})$, $q_{syj} = e^{\eta_{syj}} / \sum_{l=1}^p e^{\eta_{syl}}$ for $j = 1, \dots, p$, with linear equation

$$\eta_{syj} = \alpha_{0j} + \alpha_{sj} + y(\varphi_{0j} + \varphi_{sj}). \quad (3)$$

Politician-specific terms are set to zero for emoticon tweets (which are not associated with a specific politician), say $s = e$, such that $\eta_{eyj} = \alpha_{0j} + y\varphi_{0j}$ as a generic sentiment model. Thus all text is centered on main effects in α_0 and φ_0 , while interaction terms α_s and φ_s are identified only through their corresponding turk-scored political sentiment sample.

Results in Taddy (2013) show that $\mathbf{x}'[\boldsymbol{\varphi}_0, \boldsymbol{\varphi}_s]$ is sufficient for sentiment when \mathbf{x} is drawn from the collapsed count model implied by (3). Thus following the same logic behind our univariate SR scores in (2), $\mathbf{z}_i = [\mathbf{z}_{i0}, \mathbf{z}_{is}] = \mathbf{f}_i'[\boldsymbol{\varphi}_0, \boldsymbol{\varphi}_s]$ is a bivariate SR score for tweet i on politician s . The forward model is again proportional-odds logistic regression,

$$p(y_i \leq c) = 1/(1 + \exp[\beta_0 z_{i0} + \beta_s z_{is} - \gamma_c]), \quad (4)$$

with main β_0 and subject β_s effects. Note the absence of subject-specific γ_{sc} : a tweet containing no significant tokens (such that $z_{i0} = z_{is} = 0$) is assigned probabilities according to the overall aggregation of tweets. Such “empty” tweets have $p(-1) = 0.25$, $p(0) = 0.65$, and $p(1) = 0.1$ in the fitted model of Section 5 and are thus all classified as “neutral.”

2.3 Notes on MNIR Estimation

Estimation of MNIR models such as those in (1) and (3) follows exactly the procedures of Taddy (2013), and the

interested reader should look there for detail. Briefly, we apply the *gamma lasso* estimation algorithm, which corresponds to MAP estimation under a hierarchical gamma-Laplace coefficient prior scheme. Thus, and this is especially important for the interaction models of Section 2.1, parameters are estimated as exactly zero until a large amount of evidence has accumulated. Optimization proceeds through coordinate descent and, along with the obvious efficiency derived from collapsing observations, allows for estimation of single-factor SR models with hundreds of thousands of tokens in mere seconds. The more complicated interaction model in (3) can be estimated in less than 10 min.

To restate the MNIR strategy, we are using a simple but very high-dimensional (collapsed count) model to obtain a useful but imperfect text summary for application in low-dimensional sentiment regression. MNIR works because the multinomial is a useful representation for token counts, and this model assumption increases efficiency by introducing a large amount of information about the functional relationship between text and sentiment into the prediction problem. Implicit here is an assumption that ad-hoc forward regression can compensate for misapplication of population-average summary projections to individual document counts. Taddy (2013) presented empirical evidence that this holds true in practice, with MNIR yielding higher quality prediction at lower computational cost when compared with a variety of text regression techniques. However the design algorithms of this article are not specific to MNIR and can be combined with any sentiment prediction routine.

3. TOPIC-OPTIMAL DESIGN

Recall the introduction's pool-based design problem: choosing from the full sample of 2.1 million political tweets a subset to be scored, on mechanical turk, as either negative, neutral, or positive about the relevant politician.

A short review of some relevant literature on active learning and experimental design is in the Appendix. In our specific situation of a very high-dimensional input space (i.e., a large vocabulary), effective experimental design is tough to implement. Space-filling is impractical since limited sampling will always leave a large distance between observations. Boundary selection—where documents with roughly equal sentiment-class probabilities are selected for scoring—leads to samples that are very sensitive to model fit and is impossible in early sampling where the meaning of most terms is unknown (such that the vast majority of documents lie on this boundary). Moreover, one-at-a-time point selection implies sequential algorithms that scale poorly for large applications, while more elaborate active learning routines that solve for optimal batches of new points tend to have their own computational limits in high dimension. Finally, parameter and predictive uncertainty—which are relied upon in many active learning routines—are difficult to quantify in complicated text regression models; this includes MNIR, in which the posterior is nonsmooth and is accompanied by an ad-hoc forward regression step. The vocabulary is also growing with sample size and a full accounting of uncertainty about sentiment in unscored texts would depend heavily on a prior model for the meaning of previously unobserved words.

While the above issues make tweet selection difficult, we do have an advantage that can be leveraged in application: a huge pool of unscored documents. Our solution for text sampling is thus to look at space-filling or optimal design criteria (e.g., D-optimality) but on a reduced dimension factor decomposition of the covariate space rather than on \mathbf{X} itself. That is, although the main goal is to learn Φ for the sentiment projections of Section 2, this cannot be done until enough documents are scored and we instead look to space-fill on an *unsupervised* factor structure that can be estimated without labeled examples. This leads to what we call *factor-optimal design*. Examples of this approach include Galvanin, Macchietto, and Bezzo (2007) and Zhang and Edgar (2008), who applied optimal design criteria on principal components, and Davy and Luz (2007), a text classification contribution that applies active learning criteria to principal components fit for word counts. The proposal here is to replace generic principal component analysis with text-appropriate topic model factorization.

3.1 Multinomial Topic Factors

A K -topic model (Blei, Ng, and Jordan 2003) represents each vector of document token counts, $\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ with total $m_i = \sum_{j=1}^p x_{ij}$, as a multinomial factor decomposition

$$\mathbf{x}_i \sim \text{MN}(\omega_{i1}\theta_1 + \dots + \omega_{iK}\theta_K, m_i), \quad (5)$$

where topics $\theta_k = [\theta_{k1}, \dots, \theta_{kp}]'$ and weights ω_i are probability vectors. Hence, each topic θ_k —a vector of probabilities over words or phrases—corresponds to factor “loadings” or “rotations” in the usual factor model literature. Documents are thus characterized through a mixed-membership weighting of topic factors and ω_i is a reduced dimension summary for \mathbf{x}_i .

Briefly, this approach assumes independent prior distributions for each probability vector,

$$\begin{aligned} \omega_i &\stackrel{\text{iid}}{\sim} \text{Dir}(1/K), \quad i = 1, \dots, n, \quad \text{and} \\ \theta_k &\stackrel{\text{iid}}{\sim} \text{Dir}(1/(Kp)), \quad k = 1, \dots, K, \end{aligned} \quad (6)$$

where $\theta \sim \text{Dir}(\alpha)$ indicates a Dirichlet distribution with concentration parameter α and density proportional to $\prod_{j=1}^{\dim(\theta)} \theta_j^{\alpha}$. These $\alpha < 1$ specifications encourage a few dominant categories among mostly tiny probabilities by placing weight at the edges of the simplex. The particular specification in (6) is chosen so that prior weight, measured as the sum of concentration parameters multiplied by the dimension of their respective Dirichlet distribution, is constant in both K and p (although not in n). The model is estimated through posterior maximization (MAP estimation) as in Taddy (2012), and we employ a Laplace approximation for simulation from the conditional posterior for Ω given $\Theta = [\theta_1, \dots, \theta_K]$. The same posterior approximation allows us to estimate Bayes factors for potential values of K , and we use this to *infer* the number of topics from the data. Details are in Appendix A.2.

3.2 Topic D-Optimal Design

As a general practice, one can look to implement any space filling design in the K dimensional ω -space. For the current study, we focus on D-optimal design rules that seek to maximize the determinant of the information matrix for linear regression;

the result is thus loosely optimal under the assumption that sentiment has a linear trend in this representative factor space. The algorithm tends to select observations that are at the edges of the topic space. An alternative option that may be more robust to sentiment-topic nonlinearity is to use a latin hypercube design; this will lead to a sample that is spread evenly throughout the topic space.

In detail, we seek to select a design of documents $\{i_1, \dots, i_T\} \subset \{1, \dots, n\}$ to maximize the topic information determinant $D_T = |\mathbf{\Omega}'_T \mathbf{\Omega}_T|$, where $\mathbf{\Omega}_T = [\omega_1, \dots, \omega_T]'$ and ω_i are topic weights associated with document i_i . Since construction of exact D-optimal designs is difficult and the algorithms are generally slow (see Atkinson and Donev 1992, for an overview of both exact and approximate optimal design), we use a simple greedy search to obtain an *ordered* list of documents for evaluation in a near-optimal design.

Given $D_t = |\mathbf{\Omega}'_t \mathbf{\Omega}_t|$ for a current sample of size t , the topic information determinant after adding i_{t+1} as an additional observation is

$$D_{t+1} = |\mathbf{\Omega}'_t \mathbf{\Omega}_t + \omega'_{i_{t+1}} \omega_{i_{t+1}}| = D_t(1 + \omega'_{i_{t+1}}(\mathbf{\Omega}'_t \mathbf{\Omega}_t)^{-1} \omega_{i_{t+1}}), \quad (7)$$

due to a standard linear algebra identity. This implies that, given $\mathbf{\Omega}_t$ as the topic matrix for your currently evaluated documents, D_{t+1} is maximized simply by choosing i_{t+1} such that

$$\omega_{i_{t+1}} = \arg \max_{\omega \in \mathbf{\Omega}_t} \omega'(\mathbf{\Omega}'_t \mathbf{\Omega}_t)^{-1} \omega. \quad (8)$$

Since the topic weights are a low (K)-dimensional summary, the necessary inversion $(\mathbf{\Omega}'_t \mathbf{\Omega}_t)^{-1}$ is on a small $K \times K$ matrix and will not strain computing resources. This inverted matrix provides an operator that can quickly be applied to the pool of candidate documents (in parallel if desired), yielding a simple score for each that represents the proportion by which its inclusion increases our information determinant.

For the recursive equation in (8) to apply, the design must be initially seeded with at least K documents, such that $\mathbf{\Omega}'_t \mathbf{\Omega}_t$ will be nonsingular. We do this by starting from a simple random sample of the first $t = K$ documents (alternatively, one could use more principled space-filling in factor space, such as a latin hypercube sample). Note that again topic-model dimension reduction is crucial: for our greedy algorithm to work in the full p -dimensional token space, we would need to sample p documents before having an invertible information matrix. Since this would typically be a larger number of documents than desired for the full sample, such an approach would never move beyond the seeding stage.

In the execution of this design algorithm, the topic weights for each document must be estimated. In what we label MAP topic D-optimal design, each ω_i for document i is fixed at its MAP estimate as described in Section 3.1. As an alternative, we also consider a *marginal* topic D-optimality wherein a set of topic weights $\{\omega_{i_1}, \dots, \omega_{i_B}\}$ are sampled for each document from the approximate posterior in Appendix A.1, such that recursively D-optimal documents are chosen to maximize the *average* determinant multiplier over this set. Thus, instead of (8), marginal D-optimal i_{t+1} is selected to maximize $\frac{1}{B} \sum_b \omega'_{i_{t+1}b} (\mathbf{\Omega}'_t \mathbf{\Omega}_t)^{-1} \omega_{i_{t+1}b}$.

3.3 Note on the Domain of Factorization

The basic theme of this design framework is straightforward: fit an unsupervised factor model for \mathbf{X} and use an optimal design rule in the resulting factor space. Given a single sentiment variable, as in examples of Section 4, the \mathbf{X} to be factorized is simply the entire text corpus.

Our political twitter case study introduces the added variable of “politician,” and it is no longer clear that a single shared factorization of all tweets is appropriate. Indeed, the interaction model of Section 2.2 includes parameters (the α_{sj} and φ_{sj}) that are only identified by tweets on the corresponding politician. Given the massive amount of existing data from emoticon tweets on the other model parameters, any parameter learning from new sampling will be concentrated on these interaction parameters. Our solution in Section 5 is to apply stratified sampling: fit independent factorizations to each politician-specific subsample of tweets and obtain D-optimal designs on each. Thus we ensure a scored sample of a chosen size for each individual politician.

4. EXAMPLE EXPERIMENT

To illustrate this design approach, we consider two simple text-sentiment examples. Both are detailed in Taddy (2012, 2013) and available in the *textir* package for R. *Congress109* contains 529 legislators’ usage counts for each of 1000 phrases in the 109th U.S. Congress, and we consider party membership as the “sentiment” of interest: $y = 1$ for Republicans and 0 otherwise (two independents caucused with Democrats). *We8there* consists of counts for 2804 bigrams in 6175 online restaurant reviews, accompanied by restaurant *overall* rating on a scale of 1–5. To mimic the motivating application, we group review sentiment as negative ($y = -1$) for ratings of 1–2, neutral ($y = 0$) for 3–4, and positive ($y = 1$) for 5 (average rating is 3.95, and the full 5-class analysis is in Taddy 2013). Sentiment prediction follows the single-factor MNIR procedure of Section 2, with binary logistic forward regression $\mathbb{E}[y_i] = \exp[\gamma + \beta z_i] / (1 + \exp[\gamma + \beta z_i])$ for the congress data and proportional-odds logistic regression $p(y_i \leq c) = \exp[\gamma_c - \beta z_i] / (1 + \exp[\gamma_c - \beta z_i])$, $c = -1, 0, 1$ for the we8there data.

We fit $K = 12$ and 20 topics, respectively, to the congress109 and we8there document sets. In each case, the number of topics is chosen to maximize the approximate marginal data likelihood, as detailed in the Appendix and in Taddy (2012). Ordered sample designs were then selected following the algorithms of Section 3.2: for MAP D-optimal, using MAP topic weight estimates, and for marginal D-optimal, based upon approximate posterior samples of 50 topic weights for each document. We also consider principal component D-optimal designs, built following the same algorithm but with topic weights replaced by the same number (12 or 20) of principal components directions fit on token frequencies $\mathbf{f}_i = \mathbf{x}_i / m_i$. Finally, simple random sampling is included as a baseline and was used to seed each D-optimal algorithm with its first K observations. Each random design algorithm was repeated 100 times.

Results are shown in Figure 2, with average error rates (misclassification for congress109 and mean absolute error for we8there) reported for maximum probability classification over the entire dataset. The MAP D-optimal designs perform better




Figure 2. Average error rates on 100 repeated designs for the 109th congress and we8there examples. “MAP” is D-optimal search on MAP estimated topics, “Bayes” is our search for marginal D-optimality when sampling from the topic posterior, “PCA” is the same D-optimal search in principal components factor space, and “random” is simple random sampling. Errors are evaluated over the entire dataset. The online version of this figure is in color.

than simple random sampling, in the sense that they provide faster reduction in error rates with increasing sample size. The biggest improvements are in early sampling and error rates converge as we train on a larger proportion of the data. There is no advantage gained from using a principal component (rather than topic) D-optimal design, illustrating that misspecification of factor models can impair or eliminate their usefulness in dimension reduction. Furthermore, we were surprised to find that, in contrast with some previous studies on active learning (e.g., Taddy, Gramacy, and Polson 2011), averaging over posterior uncertainty did not improve performance: the MAP D-optimal design does as well or better than the marginal alternative, which is even outperformed by random sampling in the we8there example. Our hypothesis is that, since conditioning on Θ removes dependence across documents, sampling introduces Monte Carlo variance without providing any beneficial information about correlation in posterior uncertainty. Certainly, given that the marginal algorithm is also much more time consuming (with every operation executed B times in addition to the basic cost of sampling), it seems reasonable to focus on the MAP algorithm in application.

5. ANALYSIS OF POLITICAL SENTIMENT IN TWEETS

This section describes selection of tweets for sentiment scoring from the political Twitter data described in Section 1.1, under the design principles outlined above, along with an MNIR analysis of the results and sentiment prediction over the full collection.

5.1 Topic Factorization and D-Optimal Design

As the first step in experimental design, we apply the topic factorization of Section 3.1 independently to each politician’s tweet set. Using the Bayes factor approach of Taddy (2012), we tested K of 10, 20, 30, and 40 for each collection and, in every case, selected the simple $K = 10$ model as most probable. Although this is a smaller topic model than often seen in the literature, we have found that posterior evidence tends to favor

such simple models in corpora with short documents (see Taddy 2012, for discussion of information increase with m_i).

Across politicians, the most heavily used topic (accounting for about 20% of words in each case) always had **com**, **http**, and **via** among the top five tokens by topic lift—the probability of a token within a topic over its overall usage proportion. Hence, these topics appear to represent a Twitter-specific list of stopwords. The other topics are a mix of opinion, news, or user-specific language. For example, in the Gingrich factorization one topic accounting for 8% of text with top tokens **herman**, **cain**, and **endors** is focused on Herman Cain’s endorsement, **#teaparty** is a top token in an 8% topic that appears to contain language used by self-identified members of the Tea Party movement (this term loads heavily in a single topic for each politician we tracked), while another topic with **@danecook** as the top term accounts for 10% of traffic and is dominated by posts of unfavorable jokes and links about Gingrich by the comedian Dane Cook (and forward, or “retweets,” of these jokes by his followers).

Viewing the sentiment collection problem through these interpreted topics can be useful: since a D-optimal design looks (roughly) for large variance in topic weights, it can be seen as favoring tweets on single topics (e.g., the Cain endorsement) or rare combinations of topics (e.g., a Tea Partier retweeting a Dane Cook joke). As a large proportion of our data are retweets (near 40%), scoring those sourced from a single influential poster can yield a large reduction in predictive variance, and tweets containing contradictory topics help resolve the relative weighting of words. In the end, however, it is good to remember that the topics do not correspond to subjects in the common understanding but are simply loadings in a multinomial factor model. The experimental design described in the next section treats the fitted topics as such.

5.2 Experimental Design and Sentiment Collection

Using the MAP topic D-optimal algorithm of Section 3.2, applied to each politician’s topic factorization, we built ordered lists of tweets to be scored on Mechanical Turk: 500 for each Republican primary candidate and 750 for Obama. Worker

Table 1. Top 10 politician-specific token loadings φ_{sj} by their absolute value in MNIR

Obama		Romney		Santorum		Gingrich		Paul	
republican	15.5	fu	-10	@addthi	-11.5	bold	10.6	#p2	11.1
gop	-13.2	100%	-9.6	@newtingrich	-9.9	mash	-10	#teaparti	11
#teaparti	-12.7	lover	-9.4	clown	-9.4	ap	9.9	ht	10
#tlot	-11.9	quot	-9.4	@youtub	-9.2	obama	9.9	airplan	9.6
economi	11	anytim	-9.2	polit	-8.7	campaign	-9.9	legal	-9.5
cancer	10	abt	-8.6	speech	-8.6	lesbian	-9.7	paypal	7.4
cure	9.6	lip	-8.5	opportun	-8.2	pre	9.5	flight	6.9
ignor	9.2	incom	-8.4	disgust	-8.2	bid	9.5	rep	6.7
wors	9.2	januari	8.1	threw	-7.4	recip	-9.2	everyth	-6.4
campaign	9.2	edg	8	cultur	-7.3	america	9.1	debat	6

agreement rates varied from 78% for Obama to 85% for Paul, leading to sample sizes of 406 for Romney, 409 for Santorum, 418 for Gingrich, 423 for Paul, and 583 for Obama.

Unlike the experiment of Section 4, we have no ground truth for evaluating model performance across samples without having to pay for a large amount of Turk scoring. Instead, we propose two metrics: the number of nonzero politician-specific loadings φ_{js} , and the average entropy $-\sum_{c=-1,0,1} p_c \log(p_c)$ across tweets for each politician, where $p_c = p(y=c)$ is based on the forward proportional-odds regression described in Section 5.3. We prefer the former for measuring the *amount of sample evidence*—the number of tokens estimated as significant for politician-specific sentiment in gamma-lasso penalized estimation—as a standard statistical goal in design of experiments, but the latter corresponds to the more common machine-learning metric of classification precision (indeed, entropy calculations inform many of the close-to-boundary active learning criteria in Appendix A.1).

Results are shown in Figure 3 for the sequential addition of scored tweets from the design-ranked Turk results (sentiment regression results are deferred until Section 5.3). On the left, we see that there is a steady climb in the number of nonzero politician-specific loadings as the sample sizes increase. Although the curves flatten with more sampling, it does appear that had we continued spending money on sending tweets

to the Turk, it would have led to larger politician-sentiment dictionaries. The right plot shows a familiar pattern of early overfit (i.e., underestimated classification variance) before the mean entropy begins a slower steady decline from $t = 200$ onward.

5.3 MNIR for Subject-Specific Sentiment Analysis

After all Turk results are incorporated, we are left with 2242 scored political tweets, plus the 1.6 million emoticon tweets, and a 5566 token vocabulary. These data were used to fit the politician-interaction MNIR model detailed in Section 2.2.

The top 10 politician-specific loadings (φ_{sj}) by absolute value are shown in Table 1 (recall that these are the effect on log odds for a unit increase in sentiment; thus, e.g., negatively loaded terms occur more frequently in negative tweets). This small sample shows some large coefficients, corresponding to indicators for users or groups, news sources and events, and various other labels. For example, the Obama column results suggest that his detractors prefer to use “GOP” as shorthand for the Republican party, while his supporters simply use “republican.” However, one should be cautious about interpretation: these coefficients correspond to the partial effects of sentiment on the usage proportion for a term *given* corresponding change in relative frequency for all other terms. Moreover, these are only




Figure 3. Learning under the MAP topic D-optimal design. For increasing numbers of scored tweets added from the ranked design, the left panel shows the number of significant (nonzero) loadings in the direction of politician-specific sentiment and the right panel shows mean entropy $-\sum p_c \log(p_c)$ over the full sample. The online version of this figure is in color.

Table 2. Full sample summary statistics for politician-specific sufficient reduction scores

	$\text{cor}(z_s, z_0)$	$\text{sd}(z_s)$	\bar{z}_s
obama	-0.12	0.31	0.1
romney	0.17	0.23	-0.07
santorum	0.16	0.19	-0.19
gingrichgingrich	-0.07	0.26	-0.06
paul	0.07	0.16	0.1
emoticons	—	0.06	0.006

estimates of average correlation; this analysis is not intended to provide a causal or long-term text-sentiment model.

Summary statistics for fitted SR scores are shown in Table 2. Although we are not strictly forcing orthogonality on the factor directions— z_0 and z_s , say the emotional and political sentiment directions, respectively—the political scores have only weak correlation (absolute value < 0.2) with the generic emotional scores. This is due to an MNIR setup that estimates politician-specific loadings φ_{sj} as the sentiment effect on language about a given politician *after* controlling for generic sentiment effects. Notice that there is greater variance in political scores than in emotional scores; this is due to a few large token loadings that arise by identifying particular tweets (that are heavily retweeted) or users that are strongly associated with positive or negative sentiment. However, since we have far fewer scored political tweets than there are emoticon tweets, fewer token loadings are nonzero in the politician-specific directions than in the generic direction: φ_0 is only 7% sparse, while the φ_s are an average of 97% sparse.

Figure 4 shows fitted values in forward proportional-odds logistic regression for these SR scores. We observe some very high fitted probabilities for both true positive and negative tweets, indicating again that the analysis is able to identify a subset of similar tweets with easy sentiment classification. Tweet categorization as neutral corresponds to an absence of evidence in either direction, and neutral tweets have fitted $p(0)$ with mean around 0.6. In other applications, we have found that a large number of “junk” tweets (e.g., selling an unrelated product) requires nonproportional-odds modeling to obtain high fitted neutral probabilities, but there appears to be little junk in the current sample. As an aside, we have experimented with adding “junk” as a fourth possible categorization on Mechanical Turk but have been unable to find a presentation that avoids workers consistently getting confused between this and “neutral.”

The forward parameters are MAP estimated, using the arm package for R (Gelman et al. 2012), under diffuse t -distribution




Figure 4. In-sample sentiment fit: the forward model probabilities for each observation's true category. The online version of this figure is in color.

priors; these estimates are printed in Table 3, along with some summary statistics for the implied effect on the odds of a tweet being at or above any given sentiment level. The middle row of this table contains the average effect on log-odds for each SR score: for example, we see that Santorum tweet log-odds drop by an average of -1.1 ($e^{-1.1} \approx 0.3$) when you include his politician-specific tweet information. The bottom row shows implied effect on sentiment odds scaled for a standard deviation increase in each SR score direction: an extra deviation in emotional z_0 multiplies the odds by $e^{0.5} \approx 1.6$, while a standard deviation increase in political SR scores implies more dramatic odds multipliers of 3 (Santorum) to 8 (Gingrich). This agrees with the fitted probabilities of Figure 4 and again indicates that political directions are identifying particular users or labels and not “subjective language” in the general sense.

Figure 5 shows predicted sentiment classification for each of our 2.1 million collected political tweets, aggregated by day for each politician subject. In each case, the majority of traffic lacks enough evidence in either direction and is classified as neutral. However, some clear patterns do arise. The three “mainstream” Republicans (Romney, Santorum, and Gingrich) have far more negative than positive tweets, with Rick Santorum performing worst. Libertarian Ron Paul appears to be relatively popular on Twitter, while President Obama is the only other politician to receive (slightly) more positive than negative traffic. It is also possible to match sentiment classification changes to events; for example, Santorum's negative spike around February 20 comes after a weekend of new aggressive speeches in which he referenced Obama's “phony theology” and compared schools to “factories,” among other lines that generated controversy.

Finally, note for comparison that without the interaction terms (i.e., with only *score* as a covariate in inverse regression), the resulting univariate SR projection is dominated by emoticon-scored text. These projections turn out to be a poor summary

Table 3. MAP estimated parameters and the conditional standard deviation (ignoring variability in z) in the forward proportional-odds logistic regression $p(y_i \leq c) = (1 + \exp[\beta_0 z_{i0} + \beta_s z_{is} - \gamma_c])^{-1}$, followed by the average effect on log-odds for each SR score and exponentiated coefficients scaled according to the corresponding full-sample score standard deviation

	Intercepts γ_c		SR score coefficients β_0, β_s					
	≤ -1	≤ 0	emoticons	obama	romney	santorum	gingrich	paul
Estimate	-1.1 (0.1)	2.2 (0.1)	8.3 (1.1)	4.9 (0.5)	5.6 (0.5)	5.8 (0.5)	7.9 (1.0)	11.9 (1.1)
$\beta \times \bar{z}_s$			0.0	0.5	-0.4	-1.1	-0.5	1.2
$\exp[\beta \times \text{sd}(z)]$			1.6	4.5	3.6	2.9	7.7	6.4




Figure 5. Twitter sentiment regression full-sample predictions. Daily tweet count percentages by sentiment classification are shown with green for positive, gray neutral, and red negative.

of sentiment in the political tweets: there is little discrimination between SR scores across sentiment classes, and the in-sample misclassification rate jumps to 42% (from 13% for the model that uses politician-specific intercepts). Fitted class probabilities are little different from overall class proportions, and with true neutral tweets being less common (at 22% of our turk-scored sample), the result is that all future tweets are unrealistically predicted as either positive or negative.

6. DISCUSSION

This article makes two simple proposals for text-sentiment analysis. First, looking to optimal design in topic factor space can be useful for choosing documents to be scored. Second, sentiment can be interacted with indicator variables in MNIR to allow subject-specific inference to complement information sharing across generic sentiment.

Both techniques deserve some caution. Topic D-optimal design ignores document length, even though longer documents can be more informative; this is not a problem for the standardized Twitter format, and did not appear to harm design for our illustrative examples, but it could be an issue in other settings. In the MNIR analysis, we have observed that subject-specific sentiment loadings (driven in estimation by small sample subsets) can be dominated by news or authors specific to the given sample. While this is not technically overfit, since it is finding persistent signals in the current time period, it indicates that one should constantly update models when using these techniques for longer-term prediction.

A general lesson from this study is that traditional statistical techniques, such as experimental design and variable interaction, will apply in new areas such as text mining when used in conjunction with careful dimension reduction. Basic statistics

principles can then be relied upon to build optimal predictive models and to assess their risk and sensitivity in application.

APPENDIX

A.1 Review: Active Learning and Optimal Design

The literature on text sampling is focused on the type of design of experiments that is referred to as *active learning* in machine learning. There are two main components: optimality and adaptation. In the first, new input locations are chosen with regard to the functional form of the regression model (and possibly current parameter fits), and in the second, data are added sequentially wherever it is most needed according to a specific design criterion. Early examples of this framework include the contributions of MacKay (1992), sampling always the new point with highest predictive response variance, and Cohn (1996), choosing new inputs to maximize the expected reduction in predictive variance.

In text analysis, the work of Tong and Koller (2001) on active learning for text classification with support vector machines has been very influential. Here, the next evaluated point should be that which minimizes the expected *version space*—the set of classification rules that imply perfect separation on the current sample (standard for support vector machines, kernel expansions of the covariate space ensure such separation is possible). Hence, the criterion is analogous to Cohn's expected predictive variance, but for an overspecified algorithm without modeled variance. Tong and Koller proposed three ways to find an approximately maximizing point, the most practical of which is labeled **simple**: choose the point closest to the separating hyperplane. This is equivalent to the algorithm of Schohn and Cohn (2000).

In general, algorithms within the large literature on active learning for text regression, and similar classification problems (e.g., image sorting), follow the same theme: define a metric that summarizes “response variability” for your given prediction technique and sequentially sample inputs which maximize this metric or its expected reduction over some predefined set. For example, Yang et al. (2009) minimized approximate expected classification loss in an algorithm nearly equivalent to simple; Liere and Tadepalli (1997) generated predictions from a “committee” of classifiers and sample points where there is disagreement about the class label; and Holub, Perona, and Burl (2008) sampled to maximize the reduction in expected entropy. Since all of these methods seek to choose points near the classification boundary, it is often desirable to augment the active learning with points from a space filling design (e.g., Hu et al. 2010). Under fully Bayesian classifier active learning, as in Taddy, Gramacy, and Polson (2011), such “exploration” is automatic through accounting for posterior uncertainty about class probabilities.

A related literature from statistics is that on *optimal design*, wherein sampling is designed to optimize some function of the (traditionally linear) regression model fit; for example, one can seek to minimize parameter variance or to maximize statistical evidence. See Atkinson and Donev (1992) for an overview. Hoi, Jin, and Lyu (2006) provided an example of optimal design in text analysis. Our approach in Section 3 centers on this optimal design literature, and in particular builds upon *D-optimal* designs (Wald 1943; St. John and Draper 1975) which, if \mathbf{X} is the sample covariate matrix, seek to maximize the determinant $|\mathbf{X}'\mathbf{X}|$ (thus minimizing the determinant of coefficient covariance for an ordinary least squares regression onto \mathbf{X}). When optimal design is applied to sequential sampling problems, its goals converge with those of active learning. The main distinction is that while active learning is usually focused on adding points one-at-a-time, sequential optimal design such as in Müller and Parmigiani (1995) optimizes batch samples.

A.2 Topic Estimation and Partial Uncertainty Quantification

Topic analysis in this article follows the MAP estimation approach of Taddy (2012), yielding jointly optimal $\mathbf{\Omega}$ and $\mathbf{\Theta}$. Briefly, parameters are fit to maximize the joint posterior $L(\mathbf{\Omega}, \mathbf{\Theta})$ for $\mathbf{\Omega}$ and $\mathbf{\Theta}$ after transformed into their natural exponential family parameterization. This is equivalent to posterior maximization after adding 1 to each α prior concentration parameter and is useful for providing algorithm stability and avoiding boundary solutions.

For posterior approximation, it is also convenient to work with document topic weights (i.e., factors) transformed to this natural exponential family parameterization. That is, $\mathbf{\Omega}$ is replaced by $\mathbf{\Lambda} = \{\lambda_1, \dots, \lambda_n\}$ where for each document $\omega_k = \exp[\lambda_{k-1}] / \sum_{h=0}^{K-1} \exp[\lambda_h]$, $k = 1, \dots, K$, with the fixed element $\lambda_0 = 0$. Given MAP $\hat{\mathbf{\Lambda}}$ and $\hat{\mathbf{\Theta}}$, a Laplace approximation (e.g., Tierney and Kadane 1986) to the posterior is available as $p(\mathbf{\Lambda}, \mathbf{\Theta} | \mathbf{X}) \approx N([\hat{\mathbf{\Lambda}}, \hat{\mathbf{\Theta}}], \mathbf{H}^{-1})$, where \mathbf{H} is the log posterior Hessian (i.e., the posterior information matrix) evaluated at these MAP estimates. A further approximation replaces \mathbf{H} with its

block-diagonal, ignoring off-diagonal elements $\partial^2 L / \partial \theta_{jk} \partial \lambda_{ih}$. This allows us to avoid evaluating and inverting the full matrix \mathbf{H} , a task that is computationally impractical in large document collections. The Laplace approximation implies marginal likelihood estimates for a given K , and these are used throughout this article as the basis for selecting the number of topics. The approximate posterior also allows for topic uncertainty quantification through sampling from the *conditional* posterior for $\mathbf{\Lambda}$ (or $\mathbf{\Omega}$) given $\mathbf{\Theta}$,

$$p(\lambda_i | \mathbf{\Theta}, \mathbf{x}_i) \approx N(\hat{\lambda}_i, \mathbf{H}_i^{-1}), \quad (\text{A.1})$$

where \mathbf{H}_i has j th-row, k th-column element $h_{ijk} = \partial^2 L / \partial \lambda_{ij} \partial \lambda_{ik} = \mathbb{1}_{[j=k]} \omega_k - \omega_{ij} \omega_{ik}$. Note that document factors are independent from each other conditional on $\mathbf{\Theta}$. Our approach to posterior approximation is thus to draw $\lambda_{i1}, \dots, \lambda_{iB}$ from (A.1) and apply the logit transform to obtain $\omega_{i1}, \dots, \omega_{iB}$ as a sample from $p(\omega_i | \mathbf{\Theta}, \mathbf{x}_i)$. Although ignoring uncertainty about $\mathbf{\Theta}$ leads to only a partial assessment of variability, correlation between individual ω_i and $\mathbf{\Theta}$ decreases with n and the simple normal approximation allows fast posterior sampling.

SUPPLEMENTARY MATERIAL

The *textir* package for R is available at www.cran.r-project.org/web/packages/textir. In addition, R-code and data have been made available as supplementary material for this article.

pols.sentiment.csv: Tweet sentiment along with subject and twitter id. (csv file)

pols.tokencounts.csv: Counts for term tokens in the above tweets. (csv file)

dopt.R: Code for the various greedy design algorithms. (R-code)

congress.R: Code for the Congressional speech experimental design. (R-code)

we8there.R: Code for the we8there reviews experimental design. (R-code)

pols.mnir.R: Code for sentiment analysis based on the scored tweets. (R-code)

ACKNOWLEDGMENT

The author thanks the editor and reviewers, and participants in the Department of Energy CoDA workshop for much helpful discussion and advice.

[Received June 2012. Revised January 2013.]

REFERENCES

- Atkinson, A. C., and Donev, A. N. (1992), *Optimum Experiment Design*, Oxford: Clarendon Press. [419,424]
- Blei, D. M., Ng, A. Y., and Jordan, M. I. (2003), “Latent Dirichlet Allocation,” *Journal of Machine Learning Research*, 3, 993–1022. [418]
- Cohn, D. A. (1996), “Neural Network Exploration Using Optimal Experimental Design,” in *Neural Information Processing Systems*, pp. 679–686. [423]
- Davy, M., and Luz, S. (2007), “Dimensionality Reduction for Active Learning With Nearest Neighbour Classifier in Text Categorisation Problems,” in

- International Conference on Machine Learning and Applications*, pp. 1–8. [418]
- Galvanin, F., Macchietto, S., and Bezzo, F. (2007), “Model-Based Design of Parallel Experiments,” *Industrial & Engineering Chemistry Research*, 46, 871–882. [418]
- Gelman, A., Su, Y.-S., Yajima, M., Hill, J., Pittau, M. G., Kerman, J., and Zheng, T. (2012), *Arm: Data Analysis Using Regression and Multi-level/Hierarchical Models*, R package version 1.5-03. [422]
- Hoi, S. C. H., Jin, R., and Lyu, M. R. (2006), “Large-Scale Text Categorization by Batch Mode Active Learning,” in *International World Wide Web Conference (WWW)*, pp. 1–10. [424]
- Holub, A., Perona, P., and Burl, M. (2008), “Entropy-Based Active Learning for Object Recognition,” in *Computer Vision and Pattern Recognition*, pp. 1–8. [424]
- Hu, R., Lindstrom, P., Delany, S. J., and Namee, B. M. (2010), “Exploring the Frontier of Uncertainty Space,” in *AISTATS 2010 Poster Highlight Papers*, 1–2. [424]
- Liere, R., and Tadepalli, P. (1997), “Active Learning With Committees for Text Categorization,” in *Proceedings of the 14th National Conference on Artificial Intelligence. AAAI*, 1–6. [424]
- MacKay, D. J. C. (1992), “Information-Based Objective Functions for Active Data Selection,” *Neural Computation*, 4(4), 589–603. [423]
- Müller, P., and Parmigiani, G. (1995), “Optimal Design via Curve Fitting of Monte Carlo Experiments,” *Journal of the American Statistical Association*, 90, 1322–1330. [424]
- Porter, M. F. (1980), “An Algorithm for Suffix Stripping,” *Program*, 14, 130–137. [416]
- Schohn, G., and Cohn, D. A. (2000), “Less is More: Active Learning With Support Vector Machines,” in *Proceedings of the 17th International Conference on Machine Learning*, 1–8. [423]
- St. John, R. C., and Draper, N. R. (1975), “D-optimality for Regression Designs: A Review,” *Technometrics*, 17, 15–23. [424]
- Taddy, M. (2012), “On Estimation and Selection for Topic Models,” in *Proceedings of the 15th International Conference on Artificial Intelligence and Statistics (AISTATS 2012)*, pp. 1–12. [418,419,420,424]
- (2013), “Multinomial Inverse Regression for Text Analysis” (with discussion), *Journal of the American Statistical Association*, 108, 755–770. [416,417,418,419]
- Taddy, M., Gramacy, R. B., and Polson, N. G. (2011), “Dynamic Trees for Learning and Design,” *Journal of the American Statistical Association*, 106, 109–123. [420,424]
- Tierney, L., and Kadane, J. B. (1986), “Accurate Approximations for Posterior Moments and Marginal Densities,” *Journal of the American Statistical Association*, 81, 82–86. [424]
- Tong, S., and Koller, D. (2001), “Support Vector Machine Active Learning With Applications to Text Classification,” *Journal of Machine Learning Research*, 2, 45–66. [423]
- Wald, A. (1943), “On the Efficient Design of Statistical Investigations,” *The Annals of Mathematical Statistics*, 14, 134–140. [424]
- Yang, B., Sun, J.-T., Wang, T., and Chen, Z. (2009), “Effective Multi-Label Active Learning for Text Classification,” in *KDD 09*, pp. 917–926. [424]
- Zhang, Y., and Edgar, T. F. (2008), “PCA Combined Model-Based Design of Experiments (doe) Criteria for Differential and Algebraic System Parameter Estimation,” *Industrial & Engineering Chemistry Research*, 47, 7772–7783. [418]