

Master's Paper - 2014 - Nelson Auner

1. Abstract

This paper introduces a variant to existing models of multinomial regression for text analysis. Using the base model introduced by Taddy (2013), we extend the data-generating model to incorporate topics not explained by existing Metadata. In doing so, we seek to both increase the prediction accuracy over existing techniques, bridge the gap between multinomial regression and standard topic models, and investigate methods for discovering new topics in a corpus. We explore computational aspects of our approach, provide software for parallelization of the algorithm, and conclude by proposing areas of future research.

2. Introduction

Text Data

Multinomial models are a common way of modeling annotated text. Typically, textual information is grouped into documents, and each document is represented by counts of tokens. A document might be a single written text (e.g. an academic article), or a collection of works by the same author (e.g. all of the lyrics of an album by the rolling stones) A token is often a single word (called unigram) but may also be a sequence of two or more words (e.g. bi-grams, like 'good swimmer' is a bigram, or tri-grams, like 'I eat cheese'). The word components of tokens are often reduced to a root form by removing suffixes (e.g. 'illuminated', 'illumination' and 'illuminating' all become 'illuminate').

These tokens are then aggregated by document: For i in $i = 1, \dots, N$, the count vector $x_i = [x_{i1}, x_{i2}, \dots, x_{ip}]$ contains the number of occurrences of first, second, \dots p th token in the i th document, where p is the total number of unique tokens in all documents. This forms the complete count matrix X , where each x_{ij} is the number of occurrences of word j in document i .

Since the number of unique words that appear in a large number of documents can be extensive, we often restrict the number of tracked tokens, p to words that occur in at least two documents. We may also remove common tokens that add little meaning and are found in all documents (i.e. 'the' or 'of').

A trivial example of such content might be student's answers to the question "What did homework assignments involve?"

Some computation and formula proving, a lot of R code.
Problems, computation using R
Some computations and writing R code.
Proofs, problems, and programming work

After removing common words and stemming the remaining words, we might produce the following count matrix:

Document	Some	comp	formula	prov	R	code	use	problem	writ	program	work
1	1	1	1	1	1	1	0	0	0	0	0
2	0	1	0	0	1	0	1	1	0	0	0
3	1	1	0	0	1	0	0	0	1	0	0
4	0	0	0	1	0	0	0	1	0	1	1

Multinomial Model

We then model each document x_i as the realization of a multinomial distribution. That is,

$$x_i \sim MN(q_i, m_i)$$

Where q_i is the vector $[q_{i1}, \dots, q_{ip}]$ of token probabilities for document x_i and m_i is $\sum_{j=1}^p x_{ij}$, or the total number of tokens in document i

It is trivial to show that the maximum likelihood estimator of q_i is $f_i = x_i/m_i$, but by imposing a structure on q_i , we can model features of the data. The two most common techniques for creating structure are *topic models* and *metadata*

Topic Models

A topic model structure assumes that each document is created from a linear combination of K topics. Each topic $l = 1, \dots, K$ represents a distribution, or vector of probability weights $\omega_l = [\omega_{l1}, \dots, \omega_{lp}]$, over words. As a simple example, we can imagine a fitness store that primarily sells books on biking, running, and swimming. We can see that a probability distribution of these topics would have high probability weights on the terms (“pedal”, “helment”) for biking, (“stride”) for running, and (“breath”, “stroke”, “water”) for swimming. By denoting the proportion of topic l as θ_l , we can imagine each document as being generated by a linear combination of topics $\omega_1\theta_1 + \omega_2\theta_2 + \omega_3\theta_3$, described as the following data-generating process:

1. Choose $\theta = \theta_1, \dots, \theta_K$ the proportion of topics. (i.e., a book completely about swimming would have $\theta = (1, 0, 0)$, a book about triathalons might have $\theta = (1/3, 1/3, 1/3)$).
2. Choose m_i , the number of words in the document
3. For each word $j \in 1, 2, \dots, m_i$, choose topic l with probability θ_l . With the corresponding weighting vector ω_l , choose a word x_{ij}

Metadata

Text data is frequently accompanied by information, or metadata, about the text itself. For example, in academic journals, metadata on an article could include the number of times the article has been cited, and the journal in which the article has been published. When this metadata is believed to be relevant to the composition of the document, we use the generic term *sentiment*. For example, given a database of written movie reviews and final rating out of five $y \in (1, 2, 3, 4, 5)$, we might want to model the relationship between the words used in the document and the final rating. A simple model introduced by Taddy (2013) is that the contents of a document x with given metadata rating y , denoted x_y , is

$$x_y \sim MN(q_y, m_y) \text{ with } q_{yj} \sim \frac{\exp[\alpha_j + y\phi]}{\sum_{l=1}^p \exp[\alpha_l + y\phi_l]}$$

To determine the linear relationship between metadata y and count data x , we use Cook’s Inverse Regression method (2007) to reduce the dimension of x while maintaining its predictive power on y . That is, find ϕ such that $y_i - \phi'x_i \perp x_i$. This criteria is called *sufficient reduction*. We then take ϕ and use it to predict content x_i from metadata y_i . This technique is always helpful and is, in fact, necessary to avoid over-fitting in the many cases where the number of words p is greater than the number of documents N . The inverse regression metadata approach has a computational advantage over topic models in that when creating maximum likelihood scores, the word count data x_{ij} can be collapsed by metadata label, that is $x_y = \sum_{i=y} x_i$

Recent articles have proposed and implemented versions combining both metadata and topic modeling approaches.

3. Theory and Approach

Mixture models and cluster membership

We now turn our attention to the purpose of this paper, which is to implement an approach to discover and model “hidden”, or previously unspecified traits, across documents by grouping content unexplained by metadata.

As a simple motivating example, we might imagine a corpus of movie reviews written by several bloggers. After accounting for text information explained by the rating (i.e. relating a 5-star rating to ‘good plot’, etc), the remaining heterogeneity in the movie review content could be related traits of the blogger (e.g. gender, or home city) We may be interested in using predicting traits about bloggers given their movie reviews, and also in determining how movie review content changes across these traits.

Model Specification

Denoting the word count of a document as the vector x_i , we propose that words in a document are distributed as a multinomial with a log-link to related sentiment or topics. That is:

$$x_i \sim MN(q_{ij}, m_{ij}); q_{ij} = \frac{\exp(\alpha_j + y_i \phi_j + u_i \Gamma_{kj})}{\sum_{l=1}^p \exp(\alpha_l + y_i \phi_l + u_i \Gamma_{kl})}$$

where y_i , u_i are the metadata and cluster membership associated with document i , and ϕ_j and Γ_{kj} are the distortion coefficients for the respective metadata and factor membership. We use the subscript k to denote that each document x_i is considered a member of $k = 1, \dots, K$ clusters, with their own distortion vectors $\Gamma_1, \dots, \Gamma_K$

Our focus is on predicting cluster membership u_i and the corresponding probability distortion Γ_i . To do so, we initialize cluster membership using one of the three following methods:

1. Random Initialization
2. K-means on the word count data X
3. K-means on the residual of the word count data after incorporating metadata y (That is, give \hat{X} predicted document counts, clustering on $X - \hat{X}$)

Estimation of Parameters via Maximum a Posteriori

The negative log likelihood of a multinomial distribution can be written as

$$L(\alpha, \phi, \Gamma, u_i) = \sum_{i=1}^N \alpha + \phi v_i + u_i \Gamma_{kj} - m_i \log \left(\sum_{j=1}^p \exp[\alpha + \phi v_i + u_i \Gamma_{kj}] \right)$$

We specify laplace priors and gamma hyperprior on coefficients, as well as a gamma lasso penalty on coefficients $c(\Phi, \Gamma)$. This procedure leads us to minimize

$$L(\alpha, \Phi, \Gamma, u_i) + \sum_{j=1}^p (\alpha_j / \sigma_\alpha)^2 + c(\Phi, \Gamma)$$

The basic algorithm we use to fit coefficients α , ϕ , Γ and cluster memberships u is two main steps iterated until convergence:

1. Determine parameters α, ϕ, Γ by fitting a multinomial regression on $y_i | x_i, u_i$ with a gamma lasso penalty

2. For each document i , determine new cluster u_i membership as $\operatorname{argmax}_{k=1,\dots,K} [\ell(y_i, x_i, u_k | \alpha, \phi, \Gamma)]$

By alternating between the two steps, we aim to converge to optimal parameter estimates α, ϕ, Γ as well as optimal cluster membership u .

Computation

As noted by Taddy, Multinomial Regression enjoys the ability to collapse observations across levels of metadata. This attractive property is preserved even in our more complex cluster membership model.

Application

Congressional Speech

We applied this method to the congressional speech data of Moskowitz and Shapiro.

Convergence and Stability of Clusters

Of key interest is knowing whether or not the algorithm converges to the same clusters when run repeatedly on the same data, and also how this convergence is affected by the 3 proposed cluster initialization

At each step in classic two

Evaluation of cluster-hot-starting

Interpretation of results

Evaluating fit

Graphs

Conclusion

References